MINGQIAN ZHENG

mingqia2@andrew.cmu.edu eeelisa.github.io Google Scholar 734-834-8955

RESEARCH INTEREST

My research focuses on communication dynamics between humans and Large Language Models (LLMs), as well as interactions among multiple LLMs. I'm interested in how human-LLM communication patterns differ from human-human interactions, aiming to optimize these exchanges for improved human-AI collaboration. My work investigates LLM safety through refusal strategies, utility-safety trade-off evaluation**,** and multi-turn interactions, while also exploring multi-agent social simulations to understand LLMs' behavioral patterns and social reasoning capabilities in scenarios ranging from moral dilemma debates to non-cooperative interactions.

EDUCATION

Carnegie Mellon University *Sep.2024 - Present*
Ph.D. in Language and Information Technology
Coursework: Advanced Natural Language Processing, Inference Algorithms for Language Modeling, Multimodal Machine Learning, Ethics, Safety, and Social Impact in NLP and LLMs
Advisors: Carolyn Rosé, Maarten Sap

University of Michigan *Sep.2022 - May.2024*
Master in Survey and Data Science (Data Science Track), GPA: 4.0/4.0
Coursework: Natural Language Processing, Information Retrieval, Statistical Methods and Machine Learning, Machine Learning for Social Science
Advisor: David Jurgens, Yajuan Si

New York University Shanghai *Sep. 2018 - May 2022*
Bachelor of Science in Mathematics GPA: 3.7/4.0
Thesis: Modeling Short-term and Long-term Dynamics of User Intention for Missing-Not-At-Random Implicit Feedback in Recommendation Systems
Advisor: Shuyang Ling
Bachelor of Science in Data Science (Computer Science Track)
Thesis: User Intention Detection and Evaluation for Recommendation Systems
Advisors: Hongyi Wen, Oliver Marin
Coursework: Machine Learning, Algorithms, Data Structures, Partial Differential Equations, Information Visualization, Databases, Numerical Analysis, Math Modeling, Probability and Statistics

UNDER REVIEW

Imperfectly Cooperative Human-AI Interactions: Comparing the Impacts of Human and AI Attributes in Simulated and User Studies, Under Review
Myke C. Cohen*, **Mingqian Zheng***, Neel Bhandari, Hsien-Te Kao, Xuhui Zhou, Daniel Nguyen, Laura Cassani, Maarten Sap, Svitlana Volkova (* denotes equal contributions)

CONFERENCE PUBLICATIONS

Let Them Down Easy! Contextual Effects of LLM Guardrails on User Perceptions and Preferences, Findings of EMNLP 2025
Mingqian Zheng, Wenjia Hu, Patrick Zhao, Motahhare Eslami, Jena D. Hwang, Faeze Brahman, Carolyn Rose, Maarten Sap

Synthetic Socratic Debates: Examining Persona Effects on Moral Decision and Persuasion Dynamics, EMNLP 2025
Jiarui Liu, Yueqi Song*, Yunze Xiao*, **Mingqian Zheng***, Lindia Tjuatja, Jana Schaich Borg, Mona Diab, Maarten Sap (* denotes equal contributions)

When "A Helpful Assistant" Is Not Really Helpful: Personas in System Prompts Do Not Improve Performances of Large Language Models, Findings of EMNLP 2024
Mingqian Zheng, Jiaxin Pei, Lajanugen Logeswaran, Moontae Lee, David Jurgens

Causally Modeling the Linguistic and Social Factors that Predict Email Response, NAACL 2025
Yinuo Xu*, Hong Chen*, Sushrita Rakshit*, Aparna Ananthasubramaniam*, Omkar Yadav*, **Mingqian Zheng***, Michael Jiang*, Lechen Zhang*, Bowen Yi*, Kenan Alkiek*, Abraham Israeli*, Bangzhao Shu*, Hua Shen*, Jiaxin Pei*, Haotian Zhang*, Miriam Schirmer*, David Jurgens (* denotes equal contributions)

AWARDS

Rackham Graduate Student Research Grant at University of Michigan *Aug.2023*
University Honors Scholar at New York University *May.2022*
NYU Shanghai 2020 Recognition Award (Top 0.5%) academic merit and community contributions *Aug.2020*

RESEARCH EXPERIENCE

The Blablablab *May.2023 - Apr.2024*
Advisor: Prof. David Jurgens at University of Michigan
- Built a pipeline for systematic evaluation of social roles in Large Language Models' system prompts
- Created a list of 162 roles covering 6 types of interpersonal relationships and 8 types of occupations, and showed that adding interpersonal roles in prompts consistently improves the models' performance over a range of questions through extensive analysis of 3 popular LLMs and 2457 questions
- Conceived and implemented multiple role-prediction methodologies, optimizing for the most effective role assignment in response to specific queries by evaluating frequency, similarity, and perplexity parameters

Global Health and Population Project on Access to Care for Cardiometabolic Diseases *Nov.2022 - Apr.2024*
Advisor: Prof. David Flood at University of Michigan
- Independently conducted data harmonization following the World Health Organization (WHO) STEPwise approach to noncommunicable disease risk factor surveillance (STEPS) using Stata
- Co-authored The Lancet Global Health publication on predicting diabetes prevalence and progress toward four WHO Diabetes Targets at the country level using a Bayesian statistical modeling approach

Investigating User Profiles on Reddit: A Multi-Label Classification Study *Mar.2023 - May.2023*
Advisor: David Jurgens at University of Michigan - Ann Arbor
- Developed a multi-label classifier using RedDust dataset on Reddit, predicting user attributes (profession, gender, hobbies) through sentence transformer-based text embeddings
- Executed K-means algorithm for community detection on Reddit, analyzing qualitative aspects for insightful community dynamics

Child and Adolescent Data Lab *Oct.2022 - May.2023*
Advisor: Prof. Joseph Ryan at University of Michigan
- Conducted thorough data processing and analysis of surveys distributed by the State Court Administrative Offices on the topics of child welfare and juvenile justice systems

User Intention Detection and Evaluation for Recommendation Systems *Feb.2022 - May.2022*
Advisors: Prof. Hongyi Wen and Prof. Oliver Marin at NYU Shanghai
- Proposed an intention-aware recall metric along with a statistical mechanics of user intention extraction based on the MovieLens 1M dataset
- Verified the reliability of the metric in evaluating recommendation models by comparing the scores and performances of four DNN and RNN models

INTERNSHIP EXPERIENCE

Natural Language Processing Research Intern(Core R&D Group) at iFlytek *Jul.2021 - Oct.2021*
- Applied Recurrent Neural Networks (RNN), sequence labelling and supervised classification methods to implement word-level language identification in a code-switching dataset containing English and Spanish

- Trained the model to achieve accuracy 86% based on the cleaned dataset of 3332 English-Spanish multilingual tweets and the model was used as the baseline model for the future product

Overseas Marketing Intern at HappieWatch *Jun.2021 - Jul.2021*
- Conducted in-depth analysis of Instagram and YouTube influencer marketing data, leading to the successful negotiation and signing of contracts with three key opinion leaders (KOLs) for digital marketing campaigns.
- Efficiently supported digital marketing strategies through social media, analyzing Shopify and Google Ads data to deliver insightful weekly reports to stakeholders

EXTRACURRICULAR/LEADERSHIP EXPERIENCE

NYU Shanghai Youth League *Sep.2019 - May.2022*
Vice Secretary (Class of 2022)
- Organized and executed more than 10 school-level activities promoting cross-cultural exchange, building a platform for service and community engagement for the entire student body

NYU Shanghai student publication On Century Avenue *Sep.2019 - May.2020*
Association Initiator and Director of News and Chinese Committee
- Collaborated with a diverse team to establish an independent student publication, ensuring representation of the entire student body
- Served as a key leader in the initial core group of 6, instrumental in expanding the team to over 80 members, significantly enhancing on-campus presence, evidenced by a 500-follower increase on Instagram
- Led the committee in conducting interviews and reporting on-campus issues, effectively bridging communication between university administration and students

Shanghai Students' Federation *Aug.2018 - May.2020*
Group Leader of New Media Department
- Coordinated 20 members to write more than 50 articles on trending issues pertinent to Shanghai college students, encompassing news and survey reports

COMMUNITY SERVICES

Conference Reviewer	ACL Rolling Review 2024-2025, NeurIPS 2025
Book Reviewer	Synthesis Lectures on Human Language Technologies

SKILLS

Programming Languages	Python, R, Stata, Matlab, SQL, HTML/CSS
Programming Libraries	HuggingFace, transformers, Pandas, scikit-learn, NumPy, Matplotlib
Deep Learning Frameworks	PyTorch, TensorFlow, Keras
Statistical and Data Analysis Skills	Regression Analysis, Multilevel Modeling, Causal Inference