

MINGQIAN ZHENG

✉ mingqia2@andrew.cmu.edu 🌐 eeelisa.github.io 🎓 Google Scholar ☎ 734-834-8955

RESEARCH INTEREST

My research advances human-centered AI safety, evaluating systems by how people actually experience them rather than by benchmark scores alone. As AI shifts from answering prompts to interacting, negotiating, and acting on people's behalf, its impact turns on how it refuses, clarifies, and adapts in the moment. Working at the intersection of HCI, NLP, and social science, I study these dynamics to strengthen both the evaluation and mitigation of safety risks, increasingly in interactive and agentic settings.

EDUCATION

Carnegie Mellon University Sep. 2024 - Present

Ph.D. in Language and Information Technology

Coursework: Advanced Natural Language Processing, Inference Algorithms for Language Modeling, Multimodal Machine Learning, Ethics, Safety, and Social Impact in NLP and LLMs

Advisors: Carolyn Rosé, Maarten Sap

University of Michigan Sep. 2022 - May. 2024

Master in Survey and Data Science (Data Science Track),

GPA: 4.0/4.0

Coursework: Natural Language Processing, Information Retrieval, Statistical Methods and Machine Learning, Machine Learning for Social Science

Advisor: David Jurgens, Yajuan Si

New York University Shanghai Sep. 2018 - May 2022

Bachelor of Science in Mathematics

GPA: 3.7/4.0

Thesis: Modeling Short-term and Long-term Dynamics of User Intention for Missing-Not-At-Random Implicit Feedback in Recommendation Systems

Advisor: Shuyang Ling

Bachelor of Science in Data Science (Computer Science Track)

Thesis: User Intention Detection and Evaluation for Recommendation Systems

Advisors: Hongyi Wen, Oliver Marin

Coursework: Machine Learning, Algorithms, Data Structures, Partial Differential Equations, Information Visualization, Databases, Numerical Analysis, Math Modeling, Probability and Statistics

UNDER REVIEW

Useless but Safe? Benchmarking Utility Recovery with User Intent Clarification in Multi-Turn Conversations, arXiv

Mingqian Zheng, Malia Morgan, Liwei Jiang, Carolyn Rosé, Maarten Sap

CONFERENCE PUBLICATIONS

Imperfectly Cooperative Human-AI Interactions: Comparing the Impacts of Human and AI Attributes in Simulated and User Studies, Findings of ACL 2026

Myke C. Cohen, **Mingqian Zheng***, Neel Bhandari*, Hsien-Te Kao, Xuhui Zhou, Daniel Nguyen, Laura Cassani, Maarten Sap, Svitlana Volkova (* denotes equal contributions)

Let Them Down Easy! Contextual Effects of LLM Guardrails on User Perceptions and Preferences, Findings of EMNLP 2025

Mingqian Zheng, Wenjia Hu, Patrick Zhao, Motahhare Eslami, Jena D. Hwang, Faeze Brahman, Carolyn Rosé, Maarten Sap

Media coverage: [Forbes]

Synthetic Socratic Debates: Examining Persona Effects on Moral Decision and Persuasion Dynamics, EMNLP 2025

Jiarui Liu, Yueqi Song*, Yunze Xiao*, **Mingqian Zheng***, Lindia Tjuaatja, Jana Schaich Borg, Mona Diab, Maarten Sap (* denotes equal contributions)

When “A Helpful Assistant” Is Not Really Helpful: Personas in System Prompts Do Not Improve Performances of Large Language Models, Findings of EMNLP 2024

Mingqian Zheng, Jiaxin Pei, Lajanugen Logeswaran, Moontae Lee, David Jurgens

Media coverage: [The Markup]

Causally Modeling the Linguistic and Social Factors that Predict Email Response, NAACL 2025

Yinuo Xu*, Hong Chen*, Sushrita Rakshit*, Aparna Ananthasubramaniam*, Omkar Yadav*, **Mingqian Zheng***, Michael Jiang*, Lechen Zhang*, Bowen Yi*, Kenan Alkiek*, Abraham Israeli*, Bangzhao Shu*, Hua Shen*, Jiaxin Pei*, Haotian Zhang*, Miriam Schirmer*, David Jurgens (* denotes equal contributions)

AWARDS

Rackham Graduate Student Research Grant at University of Michigan	Aug.2023
University Honors Scholar at New York University	May.2022
NYU Shanghai 2020 Recognition Award (Top 0.5%) academic merit and community contributions	Aug.2020

RESEARCH EXPERIENCE

Human-Centered AI Safety Sep.2024 - Present

Advisors: Prof. Carolyn Rosé and Prof. Maarten Sap at Carnegie Mellon University

- Develop human-centered AI safety methods that evaluate safety as an interaction, combining controlled human studies (HCI), computational modeling (NLP), and theory from psychology and discourse studies
- Lead research on LLM guardrails, multi-turn safety evaluation, and non-cooperative human-AI interaction through controlled human studies and benchmark construction

The Blablablab May.2023 - Apr.2024

Advisor: Prof. David Jurgens at University of Michigan

- Conducted the first systematic study of how personas in LLM system prompts affect performance, curating 162 roles across 6 interpersonal relationships and 8 expertise domains
- Showed across 4 LLM families and 2,410 factual questions that adding personas does not improve performance over a no-persona control, and found that automatically selecting the best persona performs no better than random

Global Health and Population Project on Access to Care for Cardiometabolic Diseases Nov.2022 - Apr.2024

Advisor: Prof. David Flood at University of Michigan

- Independently conducted data harmonization following the World Health Organization (WHO) STEPwise approach to noncommunicable disease risk factor surveillance (STEPS) using Stata
- Co-authored The Lancet Global Health publication on predicting diabetes prevalence and progress toward four WHO Diabetes Targets at the country level using a Bayesian statistical modeling approach

INTERNSHIP EXPERIENCE

Research Intern at Allen Institute for Artificial Intelligence (AI2) May.2026 - Present

Open Ecosystem Team; Mentors: Maarten Sap and Pradeep Dasigi

- Working on safety alignment for general long-horizon agent training

Research Intern at iFlytek Jul.2021 - Oct.2021

Core R&D Group

- Built an RNN-based word-level language identification model for English-Spanish code-switching

PROFESSIONAL SERVICES

Conference Reviewer	ACL Rolling Review 2024–2025, NeurIPS 2025–2026, COLM 2026
Book Reviewer	Synthesis Lectures on Human Language Technologies